

# Create tables in Hive and write queries to access the data in the table

Download and install Apache Derby version 10.14.2.0:

[https://db.apache.org/derby/derby\\_downloads.html#For+Java+8+and+Higher](https://db.apache.org/derby/derby_downloads.html#For+Java+8+and+Higher)








## For Java 8 and Higher (releases which support lambda expressions)

- [10.14.2.0](#) (May 3, 2018 / SVN 1828579)
- [10.13.1.1](#) (October 25, 2016 / SVN 1766613)

Download and install Apache Hive version 3.1.3:

<https://downloads.apache.org/hive/hive-3.1.3/>

## Index of /hive/hive-3.1.3

Name	Last modified	Size	Description
 <a href="#">Parent Directory</a>		-	
 <a href="#">apache-hive-3.1.3-bin.tar.gz</a>	2022-04-08 17:42	312M	
 <a href="#">apache-hive-3.1.3-bin.tar.gz.asc</a>	2022-04-08 17:42	488	
 <a href="#">apache-hive-3.1.3-bin.tar.gz.sha256</a>	2022-04-08 17:42	95	
 <a href="#">apache-hive-3.1.3-src.tar.gz</a>	2022-04-08 17:42	25M	
 <a href="#">apache-hive-3.1.3-src.tar.gz.asc</a>	2022-04-08 17:42	488	
 <a href="#">apache-hive-3.1.3-src.tar.gz.sha256</a>	2022-04-08 17:42	95	

Add environment variables for Derby and Hive

System variables	
Variable	Value
__PSLockDownPolicy	0
ComSpec	C:\WINDOWS\system32\cmd.exe
DERBY_HOME	C:\db-derby-10.14.2.0-bin
DriverData	C:\Windows\System32\Drivers\DriverData
HIVE_BIN	C:\apache-hive-3.1.3-bin\bin
HIVE_HOME	C:\apache-hive-3.1.3-bin
HIVE_LIB	C:\apache-hive-3.1.3-bin\lib
IGCCSVC_DB	AOAAANCMnd8BFdFRiHoAwF/CI+sBAAAAANpif4fsnRFvva0soi3K...

Save these paths under 'path' variable

%DERBY_HOME%\bin
%HIVE_BIN%

Go to the Derby libraries directory (db-derby-10.14.2.0\lib) and copy all \*.jar files. Then, paste them within the Hive libraries directory.

Download the files from this link. Put hive-site.xml file to hive's conf location and replace hive's current guava file with this one in lib location. Also download the bin folder from link and replace the existing hive's bin folder.

<https://1drv.ms/f/s!ArSg3Xpur4Grmw0SDqW0g44T7HYU?e=wDsoBn>

Open PowerShell as administrator and go to Hadoop sbin directory and start hadoop services using the following commands:

- start-dfs.cmd
- start-yarn.cmd

```
PS C:\WINDOWS\system32> start-dfs.cmd
PS C:\WINDOWS\system32> start-yarn.cmd
starting yarn daemons
PS C:\WINDOWS\system32> jps
19344 ResourceManager
36800 DataNode
35268 Jps
35240 NodeManager
35864 NameNode
```

Open another PowerShell window and run the following command to open Derby:

StartNetworkServer -h 0.0.0.0

```
PS C:\WINDOWS\system32> StartNetworkServer -h 0.0.0.0
Wed Sep 11 19:34:46 IST 2024 : Security manager installed using the Basic server security policy.
Wed Sep 11 19:34:46 IST 2024 : Apache Derby Network Server - 10.14.2.0 - (1828579) started and ready to accept connections on port 1527
```

Type jps in the first Powershell window and check whether NetworkServerControl is running.

```
PS C:\WINDOWS\system32> jps
14144 Jps
19344 ResourceManager
36800 DataNode
35348 NetworkServerControl
35240 NodeManager
35864 NameNode
```

Go to Apache Hive's bin location with cd command and run the following command:

hive --service schematool -dbType derby --initSchema

```

PS C:\apache-hive-3.1.3-bin\bin> hive --service schematool -dbType derby -initSchema
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/C:/hadoop-3.3.6/share/hadoop/common/lib/slf4j-reload4j-1.7.36.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/C:/apache-hive-3.1.3-bin/lib/log4j-slf4j-impl-2.17.1.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Reload4jLoggerFactory]
2024-09-11 19:47:10,283 INFO conf.HiveConf: Found configuration file null
2024-09-11 19:47:10,582 INFO tools.HiveSchemaHelper: Metastore connection URL: jdbc:derby;;databaseName=metastore_db;create=true
Metastore connection URL: jdbc:derby;;databaseName=metastore_db;create=true
2024-09-11 19:47:10,582 INFO tools.HiveSchemaHelper: Metastore Connection Driver :
org.apache.derby.jdbc.EmbeddedDriver

```

```

Initialization script completed
schemaTool completed

```

Open Hive shell by typing:

hive

```

PS C:\apache-hive-3.1.3-bin\bin> hive
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/C:/hadoop-3.3.6/share/hadoop/common/lib/slf4j-reload4j-1.7.36.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/C:/apache-hive-3.1.3-bin/lib/log4j-slf4j-impl-2.17.1.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Reload4jLoggerFactory]
2024-09-11 19:57:04,346 INFO conf.HiveConf: Found configuration file null
2024-09-11 19:57:05,066 WARN common.LogUtils: hive-site.xml not found on CLASSPATH
Hive Session ID = bbc4bf36-07aa-4f43-8c81-4b2891151533
2024-09-11 19:57:05,097 INFO SessionState: Hive Session ID = bbc4bf36-07aa-4f43-8c81-4b2891151533

```

This will open the hive CLI.

Create a database using the query:

CREATE DATABASE mydata;

```

hive> CREATE DATABASE mydata;
2024-09-11 20:02:41,458 INFO conf.HiveConf: Using the default value passed in for log id: 4afabe9c-7448-4103-bc7d-134cde5ec882
2024-09-11 20:02:41,520 INFO ql.Driver: Compiling command(queryId=prath_20240911200241_a722e849-32e1-42d8-afbe-de2b5a0cba80): CREATE DATABASE mydata
2024-09-11 20:02:41,770 INFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
2024-09-11 20:02:41,770 INFO ql.Driver: Semantic Analysis Completed (retrial = false)
2024-09-11 20:02:41,786 INFO ql.Driver: Returning Hive schema: Schema(fieldSchemas:null, properties:null)
2024-09-11 20:02:41,786 INFO ql.Driver: Completed compiling command(queryId=prath_20240911200241_a722e849-32e1-42d8-afbe-de2b5a0cba80); Time taken: 0.282 seconds
2024-09-11 20:02:41,786 INFO rexec.ReExecDriver: Execution #1 of query

```

## SHOW DATABASES;

```
hive> show databases;
2024-09-11 20:04:26,698 INFO conf.HiveConf: Using the default value passed in for log id: 4afabe9c-7448-4103-bc7d-134cde5ec802
2024-09-11 20:04:26,698 INFO session.SessionState: Updating thread name to 4afabe9c-7448-4103-bc7d-134cde5ec802 main
2024-09-11 20:04:26,698 INFO ql.Driver: Compiling command(queryId=prath_20240911200426_120872cb-510b-4285-bd67-dd4fd49ebf01): show databases
2024-09-11 20:04:26,713 INFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
2024-09-11 20:04:26,713 INFO ql.Driver: Semantic Analysis Completed (retrial = false)
2024-09-11 20:04:26,760 INFO ql.Driver: Returning Hive schema: Schema(fieldSchemas:[FieldSchema(name:database_name, type:string, comment:from deserializer)], properties:null)
2024-09-11 20:04:26,823 INFO exec.ListSinkOperator: Initializing operator LIST_SINK[0]
2024-09-11 20:04:26,823 INFO ql.Driver: Completed compiling command(queryId=prath_20240911200426_120872cb-510b-4285-bd67-dd4fd49ebf01); Time taken: 0.125 seconds
2024-09-11 20:04:26,823 INFO reexec.ReExecDriver: Execution #1 of query
2024-09-11 20:04:26,823 INFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
2024-09-11 20:04:26,823 INFO ql.Driver: Executing command(queryId=prath_20240911200426_120872cb-510b-4285-bd67-dd4fd49ebf01): show databases
2024-09-11 20:04:26,823 INFO ql.Driver: Starting task [Stage-0:DOL] in serial mode
2024-09-11 20:04:26,823 INFO metastore.HiveMetaStore: 0: get_databases: @hive#
2024-09-11 20:04:26,823 INFO HiveMetaStore.audit: ugi=prath ip=unknown-ip-addr cmd=get_databases: @hive#
2024-09-11 20:04:26,838 INFO exec.DDLTask: results : 2
2024-09-11 20:04:26,870 INFO ql.Driver: Completed executing command(queryId=prath_20240911200426_120872cb-510b-4285-bd67-dd4fd49ebf01); Time taken: 0.047 seconds
OK
2024-09-11 20:04:26,870 INFO ql.Driver: OK
2024-09-11 20:04:26,870 INFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
2024-09-11 20:04:26,901 INFO Configuration.deprecation: mapred.input.dir is deprecated. Instead, use mapreduce.input.fileinputformat.inputdir
2024-09-11 20:04:26,901 INFO mapred.FileInputFormat: Total input files to process : 1
2024-09-11 20:04:26,932 INFO exec.ListSinkOperator: RECORDS_OUT_INTERMEDIATE:0, RECORDS_OUT_OPERATOR_LIST_SINK_0:2,
default
mydata
```

Switch to the new database using the following query:

## USE mydata;

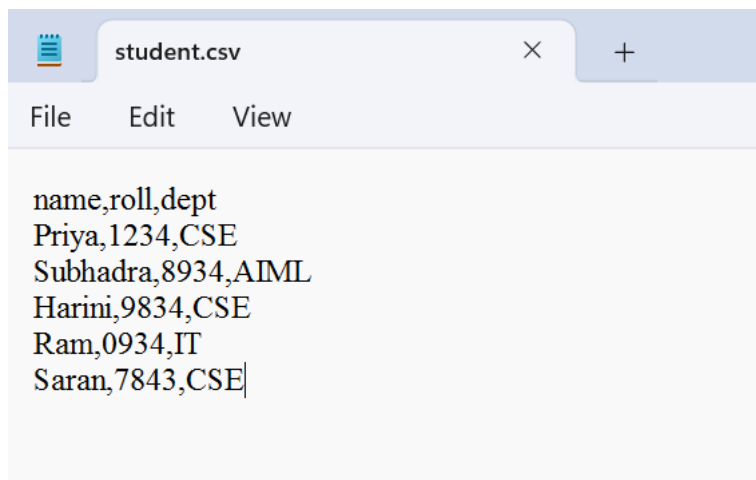
```
hive> use mydata;
2024-09-11 20:06:01,673 INFO conf.HiveConf: Using the default value passed in for log id: 4afabe9c-7448-4103-bc7d-134cde5ec802
2024-09-11 20:06:01,673 INFO session.SessionState: Updating thread name to 4afabe9c-7448-4103-bc7d-134cde5ec802 main
2024-09-11 20:06:01,675 INFO ql.Driver: Compiling command(queryId=prath_20240911200601_ac86e218-dbe4-4751-82b7-f0174b52a7a5): use mydata
2024-09-11 20:06:01,677 INFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
2024-09-11 20:06:01,677 INFO metastore.HiveMetaStore: 0: get_database: @hive#mydata
2024-09-11 20:06:01,677 INFO HiveMetaStore.audit: ugi=prath ip=unknown-ip-addr cmd=get_database: @hive#mydata
2024-09-11 20:06:01,677 INFO ql.Driver: Semantic Analysis Completed (retrial = false)
2024-09-11 20:06:01,677 INFO ql.Driver: Returning Hive schema: Schema(fieldSchemas:null, properties:null)
2024-09-11 20:06:01,677 INFO ql.Driver: Completed compiling command(queryId=prath_20240911200601_ac86e218-dbe4-4751-82b7-f0174b52a7a5); Time taken: 0.002 seconds
2024-09-11 20:06:01,677 INFO reexec.ReExecDriver: Execution #1 of query
2024-09-11 20:06:01,677 INFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
2024-09-11 20:06:01,677 INFO ql.Driver: Executing command(queryId=prath_20240911200601_ac86e218-dbe4-4751-82b7-f0174b52a7a5): use mydata
2024-09-11 20:06:01,677 INFO ql.Driver: Starting task [Stage-0:DOL] in serial mode
2024-09-11 20:06:01,677 INFO metastore.HiveMetaStore: 0: get_database: @hive#mydata
2024-09-11 20:06:01,677 INFO HiveMetaStore.audit: ugi=prath ip=unknown-ip-addr cmd=get_database: @hive#mydata
2024-09-11 20:06:01,692 INFO metastore.HiveMetaStore: 0: get_database: @hive#mydata
2024-09-11 20:06:01,692 INFO HiveMetaStore.audit: ugi=prath ip=unknown-ip-addr cmd=get_database: @hive#mydata
2024-09-11 20:06:01,692 INFO ql.Driver: Completed executing command(queryId=prath_20240911200601_ac86e218-dbe4-4751-82b7-f0174b52a7a5); Time taken: 0.015 seconds
OK
2024-09-11 20:06:01,692 INFO ql.Driver: OK
2024-09-11 20:06:01,692 INFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
Time taken: 0.018 seconds
2024-09-11 20:06:01,692 INFO CliDriver: Time taken: 0.018 seconds
2024-09-11 20:06:01,692 INFO conf.HiveConf: Using the default value passed in for log id: 4afabe9c-7448-4103-bc7d-134cde5ec802
2024-09-11 20:06:01,692 INFO session.SessionState: Resetting thread name to main
```

Create a table using the following command:

## CREATE TABLE students(name STRING, roll INT, dept STRING);

```
hive> create table students(name STRING,roll INT,dept STRING);
2024-09-11 20:08:36,936 INFO conf.HiveConf: Using the default value passed in for log id: 4afabe9c-7448-4103-bc7d-134cde5ec802
2024-09-11 20:08:36,936 INFO session.SessionState: Updating thread name to 4afabe9c-7448-4103-bc7d-134cde5ec802 main
2024-09-11 20:08:36,952 INFO ql.Driver: Compiling command(queryId=prath_20240911200836_07346f06-a91-4601-830a-a1e65abef2d3): create table students(name STRING,roll INT,
dept STRING)
2024-09-11 20:08:36,968 INFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
2024-09-11 20:08:36,983 INFO parse.CalcitePlanner: Starting Semantic Analysis
2024-09-11 20:08:36,983 INFO parse.CalcitePlanner: Creating table mydata.students position13
2024-09-11 20:08:36,999 INFO metastore.HiveMetaStore: 0: get_database: @hive#mydata
2024-09-11 20:08:36,999 INFO HiveMetaStore.audit: ugi=prath ip=unknown-ip-addr cmd=get_database: @hive#mydata
2024-09-11 20:08:36,999 INFO ql.Driver: Semantic Analysis Completed (retrial = false)
2024-09-11 20:08:37,014 INFO ql.Driver: Returning Hive schema: Schema(fieldSchemas:null, properties:null)
2024-09-11 20:08:37,014 INFO ql.Driver: Completed compiling command(queryId=prath_20240911200836_07346f06-a91-4601-830a-a1e65abef2d3); Time taken: 0.062 seconds
2024-09-11 20:08:37,014 INFO reexec.ReExecDriver: Execution #1 of query
2024-09-11 20:08:37,014 INFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
2024-09-11 20:08:37,014 INFO ql.Driver: Executing command(queryId=prath_20240911200836_07346f06-a91-4601-830a-a1e65abef2d3): create table students(name STRING,roll INT,
dept STRING)
2024-09-11 20:08:37,014 INFO ql.Driver: Starting task [Stage-0:DOL] in serial mode
2024-09-11 20:08:37,155 INFO metastore.HiveMetaStore: 0: create_table: Table(tableName=students, dbName=mydata, owner=prath, createTime=1726085517, lastAccessTime=0, ret
ention=0, sd:StorageDescriptor(cols:[FieldSchema(name=name, type:string, comment:null)], FieldSchema(name=roll, type:int, comment:null), FieldSchema(name=dept, type:string,
comment:null)], location:null, inputFormat=org.apache.hadoop.mapred.TextInputFormat, outputFormat=org.apache.hadoop.hive ql.io.HiveIgnoreKeyTextOutputFormat, compressed=false, numBuckets=1, serdeInfo:SerDeInfo(name=null, serializationLib=org.apache.hadoop.hive.serde2.lazy.LazySimpleSerDe, parameters:{serialization-format=1}), bucketCols:[], p
artitionKeys:[], parameters:{}, skewedInfo:SkewedInfo(skewedColumnNames:[], skewedColumnValues:[], skewedColumnValueLocations:[]), storedAsSubDirectories=false, partitionKeys:[], pa
rtitions:[], totalSize=0, numRows=0, rawDataSize=0, COLUMN_STATS_ACCURATE={\"BASIC STATS\"=true,\"COLUMN STATS\"={\"dept\"=true,\"name\"=true,\"roll\"=true}}, numFiles=0, bucket
version=2}, viewOriginalText=null, viewExpandedText=null, tableType=MANAGED TABLE, privileges:[PrincipalPrivilegeSet(userPrivileges:[prath:[PrivilegeGrantInfo(privilege:
ERT, createTime=-1, grantor=prath, grantorType=USER, grantOption=true), PrivilegeGrantInfo(privilege:SELECT, createTime=-1, grantor=prath, grantorType=USER, grantOption:
true), PrivilegeGrantInfo(privilege:UPDATE, createTime=-1, grantor=prath, grantorType=USER, grantOption=true), PrivilegeGrantInfo(privilege:DELETE, createTime=-1, grantor=pr
ath, grantorType=USER, grantOption=true)], groupPrivileges:[], rolePrivileges:[], temporary=false, dbName=mydata, owner=prath, createTime=1726085517, lastAccessTime=0, retention=0, sd:StorageDescriptor(cols:[FieldSchema(name=name, type:string, comment:null)], FieldSchema(name=roll, type:int, comment:null), F
ieldSchema(name=dept, type:string, comment:null)], location:null, inputFormat=org.apache.hadoop.mapred.TextInputFormat, outputFormat=org.apache.hadoop.hive ql.io.HiveIgnore
TextOutputFormat, compressed=false, numBuckets=1, serdeInfo:SerDeInfo(name=null, serializationLib=org.apache.hadoop.hive.serde2.lazy.LazySimpleSerDe, parameters:{serial
```

## Create a csv file



Add the file to HDFS using the command:

```
hdfs dfs -put C:/Semester7/DataAnalytics/Lab/Ex5/student.csv /user/hive
```

```
C:\Windows\System32>hdfs dfs -put C:/Semester7/DataAnalytics/Lab/Ex5/student.csv /user/hive  
C:\Windows\System32>
```

### File contents

name	roll	dept
Priya	1234	CSE
Subhadra	8934	AIML
Harini	9834	CSE
Ram	0934	IT
Saran	7843	CSE

Add this data to the table using the command:

```
LOAD DATA INPATH '/user/hive/student_data.csv' INTO TABLE students;
```



```
hive> LOAD DATA INPATH '/user/hive/student.csv' INTO TABLE students;
2024-09-11 20:28:18,593 INFO conf.HiveConf: Using the default value passed in for log id: 4afabe9c-7448-4103-bc7d-134cde5ec802
2024-09-11 20:28:18,594 INFO session.SessionState: Updating thread name to 4afabe9c-7448-4103-bc7d-134cde5ec802 main
2024-09-11 20:28:18,597 INFO ql.Driver: Compiling command(queryId=prath_20240911202818_c4fa0a8a-374b-4060-b98d-4a29919bb65e): LOAD DATA INPATH '/user/hive/student.csv' IN
TABLE students
2024-09-11 20:28:18,614 INFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
2024-09-11 20:28:18,620 INFO metastore.HiveMetaStore: 0: get_table : tbl=hive.mydata.students
2024-09-11 20:28:18,620 INFO HiveMetaStore.audit: ugi=prath ip=unknown-ip-addr cmd=get_table : tbl=hive.mydata.students
2024-09-11 20:28:18,874 INFO ql.Driver: Semantic Analysis Completed (retrial = false)
2024-09-11 20:28:18,874 INFO ql.Driver: Returning Hive schema: Schema(fieldSchemas:null, properties:null)
2024-09-11 20:28:18,875 INFO ql.Driver: Completed compiling command(queryId=prath_20240911202818_c4fa0a8a-374b-4060-b98d-4a29919bb65e); Time taken: 0.279 seconds
2024-09-11 20:28:18,875 INFO rexec.ReExecDriver: Execution #1 of query
2024-09-11 20:28:18,875 INFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
2024-09-11 20:28:18,876 INFO ql.Driver: Executing command(queryId=prath_20240911202818_c4fa0a8a-374b-4060-b98d-4a29919bb65e): LOAD DATA INPATH '/user/hive/student.csv' IN
TABLE students
2024-09-11 20:28:18,877 INFO ql.Driver: Starting task [Stage-0:MOVE] in serial mode
2024-09-11 20:28:18,878 INFO metastore.HiveMetaStore: 0: Cleaning up thread local RawStore...
2024-09-11 20:28:18,878 INFO HiveMetaStore.audit: ugi=prath ip=unknown-ip-addr cmd=Cleaning up thread local RawStore...
2024-09-11 20:28:18,879 INFO metastore.HiveMetaStore: 0: Done cleaning up thread local RawStore
2024-09-11 20:28:18,879 INFO HiveMetaStore.audit: ugi=prath ip=unknown-ip-addr cmd=Done cleaning up thread local RawStore
Loading data to table mydata.students
2024-09-11 20:28:18,880 INFO exec.Task: Loading data to table mydata.students from hdfs://localhost:9000/user/hive/student.csv
2024-09-11 20:28:18,883 INFO metastore.HiveMetaStore: 0: Opening raw store with implementation class:org.apache.hadoop.hive.metastore.ObjectStore
```

## SHOW TABLES;

```
hive> show tables;
2024-09-11 20:29:52,246 INFO conf.HiveConf: Using the default value passed in for log id: 4afabe9c-7448-4103-bc7d-134cde5ec802
2024-09-11 20:29:52,246 INFO session.SessionState: Updating thread name to 4afabe9c-7448-4103-bc7d-134cde5ec802 main
2024-09-11 20:29:52,249 INFO ql.Driver: Compiling command(queryId=prath_20240911202952_4ff4a7f6-94f0-4c59-85db-16b2dc73fd05): show tables
2024-09-11 20:29:52,260 INFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
2024-09-11 20:29:52,262 INFO metastore.HiveMetaStore: 0: get_database: @hive#mydata
2024-09-11 20:29:52,262 INFO HiveMetaStore.audit: ugi=prath ip=unknown-ip-addr cmd=get_database: @hive#mydata
2024-09-11 20:29:52,263 INFO ql.Driver: Semantic Analysis Completed (retrial = false)
2024-09-11 20:29:52,264 INFO ql.Driver: Returning Hive schema: Schema(fieldSchemas:[FieldSchema(name:tab_name, type:string, comment:from deserializer)], properties:null)
2024-09-11 20:29:52,265 INFO exec.ListSinkOperator: Initializing operator LIST_SINK[0]
2024-09-11 20:29:52,266 INFO ql.Driver: Completed compiling command(queryId=prath_20240911202952_4ff4a7f6-94f0-4c59-85db-16b2dc73fd05); Time taken: 0.017 seconds
2024-09-11 20:29:52,266 INFO rexec.ReExecDriver: Execution #1 of query
2024-09-11 20:29:52,266 INFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
2024-09-11 20:29:52,266 INFO ql.Driver: Executing command(queryId=prath_20240911202952_4ff4a7f6-94f0-4c59-85db-16b2dc73fd05): show tables
2024-09-11 20:29:52,267 INFO ql.Driver: Starting task [Stage-0:DOL] in serial mode
2024-09-11 20:29:52,267 INFO metastore.HiveMetaStore: 0: get_database: @hive#mydata
2024-09-11 20:29:52,267 INFO HiveMetaStore.audit: ugi=prath ip=unknown-ip-addr cmd=get_database: @hive#mydata
2024-09-11 20:29:52,268 INFO metastore.HiveMetaStore: 0: get_tables: db=@hive#mydata pat=.*
2024-09-11 20:29:52,269 INFO HiveMetaStore.audit: ugi=prath ip=unknown-ip-addr cmd=get_tables: db=@hive#mydata pat=.*
2024-09-11 20:29:52,276 INFO ql.Driver: Completed executing command(queryId=prath_20240911202952_4ff4a7f6-94f0-4c59-85db-16b2dc73fd05); Time taken: 0.01 seconds
OK
2024-09-11 20:29:52,276 INFO ql.Driver: OK
2024-09-11 20:29:52,277 INFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
2024-09-11 20:29:52,280 INFO mapred.FileInputFormat: Total input files to process : 1
2024-09-11 20:29:52,284 INFO exec.ListSinkOperator: RECORDS_OUT_INTERMEDIATE:0, RECORDS_OUT_OPERATOR_LIST_SINK_0:1,
students
Time taken: 0.029 seconds, Fetched: 1 row(s)
2024-09-11 20:29:52,288 INFO CliDriver: Time taken: 0.029 seconds, Fetched: 1 row(s)
2024-09-11 20:29:52,288 INFO conf.HiveConf: Using the default value passed in for log id: 4afabe9c-7448-4103-bc7d-134cde5ec802
2024-09-11 20:29:52,289 INFO session.SessionState: Resetting thread name to main
```

## DESC students;

```
hive> desc students;
2024-09-11 20:30:26,611 INFO conf.HiveConf: Using the default value passed in for log id: 4afabe9c-7448-4103-bc7d-134cde5ec802
2024-09-11 20:30:26,611 INFO session.SessionState: Updating thread name to 4afabe9c-7448-4103-bc7d-134cde5ec802 main
2024-09-11 20:30:26,614 INFO ql.Driver: Compiling command(queryId=prath_20240911203026_416d1533-b726-44ad-9ffc-6012fb2edab4): desc students
2024-09-11 20:30:26,625 INFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
2024-09-11 20:30:26,628 INFO metastore.HiveMetaStore: 0: get_table : tbl=hive.mydata.students
2024-09-11 20:30:26,628 INFO HiveMetaStore.audit: ugi=prath ip=unknown-ip-addr cmd=get_table : tbl=hive.mydata.students
2024-09-11 20:30:26,636 INFO parse.DDLSemanticAnalyzer: analyzeDescribeTable done
2024-09-11 20:30:26,636 INFO ql.Driver: Semantic Analysis Completed (retrial = false)
2024-09-11 20:30:26,637 INFO ql.Driver: Returning Hive schema: Schema(fieldSchemas:[FieldSchema(name:col_name, type:string, comment:from deserializer), FieldSchema(name:da
a_type, type:string, comment:from deserializer), FieldSchema(name:comment, type:string, comment:from deserializer)], properties:null)
2024-09-11 20:30:26,638 INFO exec.ListSinkOperator: Initializing operator LIST_SINK[0]
2024-09-11 20:30:26,639 INFO ql.Driver: Completed compiling command(queryId=prath_20240911203026_416d1533-b726-44ad-9ffc-6012fb2edab4); Time taken: 0.025 seconds
2024-09-11 20:30:26,639 INFO rexec.ReExecDriver: Execution #1 of query
2024-09-11 20:30:26,640 INFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
2024-09-11 20:30:26,641 INFO ql.Driver: Executing command(queryId=prath_20240911203026_416d1533-b726-44ad-9ffc-6012fb2edab4): desc students
2024-09-11 20:30:26,642 INFO ql.Driver: Starting task [Stage-0:DDL] in serial mode
2024-09-11 20:30:26,643 INFO metastore.HiveMetaStore: 0: get_table : tbl=hive.mydata.students
2024-09-11 20:30:26,643 INFO HiveMetaStore.audit: ugi=prath ip=unknown-ip-addr cmd=get_table : tbl=hive.mydata.students
2024-09-11 20:30:26,654 INFO ql.Driver: Completed executing command(queryId=prath_20240911203026_416d1533-b726-44ad-9ffc-6012fb2edab4); Time taken: 0.013 seconds
OK
2024-09-11 20:30:26,655 INFO ql.Driver: OK
2024-09-11 20:30:26,655 INFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
2024-09-11 20:30:26,659 INFO mapred.FileInputFormat: Total input files to process : 1
2024-09-11 20:30:26,663 INFO exec.ListSinkOperator: RECORDS_OUT_INTERMEDIATE:0, RECORDS_OUT_OPERATOR_LIST_SINK_0:3,
name
string
int
roll
dept
string
Time taken: 0.042 seconds, Fetched: 3 row(s)
2024-09-11 20:30:26,671 INFO CliDriver: Time taken: 0.042 seconds, Fetched: 3 row(s)
2024-09-11 20:30:26,672 INFO conf.HiveConf: Using the default value passed in for log id: 4afabe9c-7448-4103-bc7d-134cde5ec802
2024-09-11 20:30:26,673 INFO session.SessionState: Resetting thread name to main
```

Display the contents of the table using the following command:

SELECT \* FROM students;

```
hive> SELECT * FROM STUDENTS;
2024-09-11 20:31:51,244 INFO conf.HiveConf: Using the default value passed in for log id: 4afabe9c-7448-4103-bc7d-134cde5ec802
2024-09-11 20:31:51,244 INFO session.SessionState: Updating thread name to 4afabe9c-7448-4103-bc7d-134cde5ec802 main
2024-09-11 20:31:51,247 INFO ql.Driver: Compiling command(queryId=prath_20240911203151_118ce9a0-bcca-4768-8919-3962fd53b82c): SELECT * FROM STUDENTS
2024-09-11 20:31:51,258 INFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
2024-09-11 20:31:51,258 INFO parse.CalcitePlanner: Starting Semantic Analysis
2024-09-11 20:31:51,261 INFO parse.CalcitePlanner: Completed phase 1 of Semantic Analysis
2024-09-11 20:31:51,261 INFO parse.CalcitePlanner: Get metadata for source tables
2024-09-11 20:31:51,261 INFO metastore.HiveMetaStore: 0: get_table : tbl=hive.mydata.students
2024-09-11 20:31:51,261 INFO HiveMetaStore.audit: ugi=prath ip=unknown-ip-addr cmd=get_table : tbl=hive.mydata.students
2024-09-11 20:31:51,269 INFO parse.CalcitePlanner: Get metadata for subqueries
2024-09-11 20:31:51,271 INFO parse.CalcitePlanner: Get metadata for destination tables
2024-09-11 20:31:51,296 INFO ql.Context: New scratch dir is hdfs://localhost:9000/tmp/hive/prath/4afabe9c-7448-4103-bc7d-134cde5ec802/hive_2024-09-11_20-31-51_256_143077385
127262073-1
2024-09-11 20:31:51,296 INFO parse.CalcitePlanner: Completed getting MetaData in Semantic Analysis
```

```
name,roll,dept NULL NULL
Priya,1234,CSE NULL NULL
Subhadra,8934,AI ML NULL NULL
Harini,9834,CSE NULL NULL
Ram,0934,IT NULL NULL
Saran,7843,CSE NULL NULL
Time taken: 1.245 seconds, Fetched: 6 row(s)
2024-09-11 20:31:52,530 INFO CliDriver: Time taken: 1.245 seconds, Fetched: 6 row(s)
2024-09-11 20:31:52,530 INFO conf.HiveConf: Using the default value passed in for log id: 4afabe9c-7448-4103-bc7d-134cde5ec802
2024-09-11 20:31:52,537 INFO session.SessionState: Resetting thread name to main
hive>
```